When Misfortunes Change: Reversing Undemocratic Learning in the Lab

Hector Bahamonde*[1,2]

[1]Senior Researcher, INVEST Research Flagship Centre, University of Turku, Finland
[2]Project Researcher, Social Science Research Institute, Åbo Akademi University, Finland

June 18, 2026

Abstract

Democratic backsliding research has established that citizens can learn to tolerate norm violations when structural threats make such tolerance instrumentally attractive. This proposal asks whether that learning is reversible when those structural conditions improve. I implement a three-wave economic laboratory panel experiment structured as a hypothetical election campaign. Wave 1 induces a democratic threat environment in which a rule-breaking candidate can credibly promise targeted benefits. Wave 2 introduces an experimentally randomized discontinuity: structural conditions either improve sharply (treatment) or persist (control). Wave 3 culminates in an incentivized election. Voting is payoff-relevant: subjects face information acquisition costs and a voting cost to reduce cheap talk, and decision quality is evaluated using the correct-voting framework. Specifically, after the ballot, we are able to compute each subject's utility-maximizing candidate and pay an extra bonus when the vote matches that benchmark. The design is mixed: treatment assignment is between-subject, while mechanism-stage quantities are tracked within individuals across waves. The primary substantive outcome is Wave-3 democratic reversal: whether subjects who endorsed the antidemocratic option under the Wave-1 threat environment switch to the pro-democratic candidate in the final incentivized election. I process-trace the causal chain by measuring information search, institutional-capacity perceptions, and democratic-collapse risk before and after the Wave-2 discontinuity. Correct voting is used instrumentally as an incentive-compatible scoring rule: it makes information acquisition and final voting consequential by rewarding subjects when their Wave-3 vote matches their own full-information utility-maximizing choice. It is therefore part of the experimental architecture rather than a second substantive outcome.

*hector.bahamonde@utu.fi; `https://www.hectorbahamonde.com`

1 Reversing Undemocratic Learning

The most intriguing finding in recent democratic backsliding research is not that citizens tolerate norm violations (Graham and Svolik 2020). My claim is that they *learn* to do so. As citizens secularly experience norm-violating leaders and accumulate information about non-democratic behaviors (Levitsky and Ziblatt 2018), they recalibrate expectations, adjust their definitions of democracy itself (Ananda and Bol 2021), and gradually become "normalized" to authoritarianism (Grillo and Prato 2023). Yet this learning process raises a critical question: **Can that undemocratic learning be reversed?** If citizens can learn authoritarian tolerance, can they *un*-learn it when the institutional conditions and threats that gave rise to such undemocratic leaders change? And critically, which institutional conditions trigger such reversals? These questions are fundamental for democratic resilience, and importantly, for *defending* democracy (Capoccia 2007).

I implement a three-wave economic laboratory panel experiment structured as a hypothetical election campaign (see prototype here). The central purpose of using three waves is to induce, observe, and then test *reversals* in antidemocratic tolerance under changing structural conditions, while explicitly modeling the role of memory in campaign learning and vote choice. The memory literature suggests that persuasion and framing effects often decay quickly unless information remains accessible or is repeatedly refreshed, and that opinion formation can rely either on what is retrievable from memory at the moment of choice or on a more durable "online" tally (Hill et al. 2013; Lecheler and De Vreese 2011). This matters for democratic learning: if citizens rely disproportionately on *recent* information or short-term recall, then support for (anti)democratic candidates may reflect what remains *salient* in memory rather than *stable* commitments. Thus, and consistent with political-economy experiments on short-term memory, limited memory can generate polarization because what voters recently recall shapes what they infer about the "correct" policy (Levy and Razin 2025).

For the final wave ("voting day"), I will exploit the "correct voting" framework (Lau and Redlawsk 1997): a vote is *correct* if it matches the alternative the subject would select under conditions of full information given their own stated preferences (captured in a questioner on week 1; see Figure 3). In practice, I operationalize this as a utility-maximization criterion, where the correct vote is the candidate with the highest subject-specific utility score once all relevant information is revealed. Thus, the design is explicitly incentive compatible: subjects face real utility tradeoffs and explicit costs for information search and voting to reduce cheap talk (Morton and Williams 2010).

These considerations motivate a *panel* design rather than a one-shot experiment. First, "correct voting" is inherently a dynamic object in campaign contexts: information arrives *sequentially* (Redlawsk 2001), voters search selectively, and decision quality reflects the joint influence of motivation, cognitive capacity, heuristics, and task demands over time (Lau and Redlawsk 1997; Lau, Andersen, and Redlawsk 2008). Second, to evaluate *reversal*, we must observe the same individuals before and after the Wave-2 discontinuity. Third, a panel design allows us to process-trace (Gerring 2004; Levy 2008; Mahoney 2010; Collier 2011) whether changes in vote choice are driven by (i) updated learning, or (ii) simple recency and accessibility effects (i.e., what is remembered at the moment of voting). This distinction is important because learning can fail to correct even with abundant feedback when subjects operate with incorrect information (Esponda, Vespa, and Yuksel 2024). Finally, my design is mixed: the Wave-2 discontinuity is a *between-subject* treatment, while key outcomes and mechanisms are tracked *within individuals* across waves, allowing me to estimate both cross-sectional treatment effects and within-subject updating trajectories.

This structure clarifies the status of the outcome variable. The proposal has a single primary substantive outcome: democratic reversal in the Wave-3 vote. The earlier measurements are not competing outcome variables. They are process-tracing observations that allow the design to evaluate whether the hypothesized causal chain is visible before the final behavioral endpoint is reached. Wave 1 identifies the initial state of the sequence: antidemocratic endorsement under threat. Wave 2 introduces the randomized structural discontinuity and measures whether the beliefs that made antidemocratic governance instrumentally attractive are updated when structural conditions change. Wave 3 then observes whether those intermediate updates culminate in final democratic reversal. In this sense, the panel design is not a sequence of repeated endpoints, but a temporally ordered test of a causal mechanism.

2 Brief Theoretical Framework

My rationale is simple: if antidemocratic attitudes are learned under particular structural conditions, then changes in those conditions should also reshape the beliefs that sustain them. For example, Voelkel et al. (2024) find that antidemocratic attitudes can be reduced through cognitive shifts: citizens update misperceptions about opponents, revise assessments of democratic threat, and respond to cues that re-anchor democratic norms. I build on that insight by shifting the level of analysis. Rather than asking whether brief informational treatments can generate these updates, I ask whether (lab-simulated) structural reversals can generate the same updates.

Importantly, the broader literature on democratic backsliding suggests that citizens tolerate norm-eroding behavior at least partly through affective polarization mechanisms (Orhan 2022). Citizens with strong in-party attachment oppose particular constitutional protections when their party is in power, but support them when out of power—a dynamic Kingzette and colleagues term the "cue-taking mechanism" (Kingzette et al. 2021). What remains unexplored, however, is whether citizens can learn to *reverse* this tolerance when structural conditions improve.

The institutional crisis used in the experiment is therefore designed as a governance-capacity crisis rather than as a security emergency. This distinction is central. External war, terrorism, or domestic security breakdowns may activate democratic states of exception in which even constitutional democracies permit temporary restrictions on individual freedoms. Such scenarios would confound the object of the experiment: support for antidemocratic rule-breaking would be difficult to distinguish from support for legally authorized emergency powers. Instead, the Wave-1 crisis describes a prolonged failure of ordinary democratic governance: parliament has been unable to pass a budget for eight months, public-sector planning has stalled, healthcare waiting times have doubled, and economic growth has stagnated. This setting creates a credible institutional-performance problem without invoking emergency law. It also matches the empirical conditions under which democratic tradeoffs are most theoretically informative. In Finland, Saikkonen and Christensen find that even in a high-trust and low-polarization democracy, some voters tolerate authoritarian leader behavior when it is attached to policy congruence, especially on salient issues such as immigration (Saikkonen and Christensen 2023). Cross-national evidence similarly suggests that citizens attach greater weight to elections and prosperity than to horizontal constraints on the executive: citizens are more responsive to economic well-being than to civil liberties, and much less responsive to executive constraints (Neundorf et al. 2025, pp. 14–15). The crisis therefore places participants precisely at the tradeoff identified in this literature: one candidate promises slower but constitutionally constrained problem-solving through parliamentary bargaining, transparent budgeting, and judicially reviewable administrative procedures, whereas the opposing candidate proposes to temporarily centralize budgetary and administrative decisions in the executive in order to bypass gridlock, accelerate healthcare funding, and restore growth, but with reduced parliamentary oversight and weaker institutional checks. This is the most appropriate crisis for the Finnish case because it is plausible, non-security-based, and experimentally clean: it makes antidemocratic governance instrumentally attractive through institutional frustration and promised performance, while avoiding a scenario in which participants could reasonably interpret restrictions on checks and liberties as ordinary democratic emergency powers.

3 Methodological Design

Because the object of interest is reversal rather than static support for antidemocratic governance, the design observes the same construct at theoretically ordered moments in the campaign. The sequencing is central to the "process-tracing logic" (Collier 2011). The aim is to recover the causal chain rather than capturing the pre-post "effects of causes" (Pearl 2015). Process tracing is especially appropriate here because it examines whether a hypothesized causal chain is visible in the sequence of *intermediate observations* connecting treatment and outcome (Bennett and Checkel 2014). The primary behavioral outcome is Wave-3 democratic reversal. Earlier wave-specific measures are mechanism-stage observations: they record the initial endorsement of antidemocratic governance under threat, subsequent belief updating after the structural discontinuity, and the information-search patterns that connect the two. Correct voting is not treated as a separate substantive outcome. It is used instrumentally as an incentive-compatible scoring rule that makes costly information search and the final vote consequential (see Figure 1).

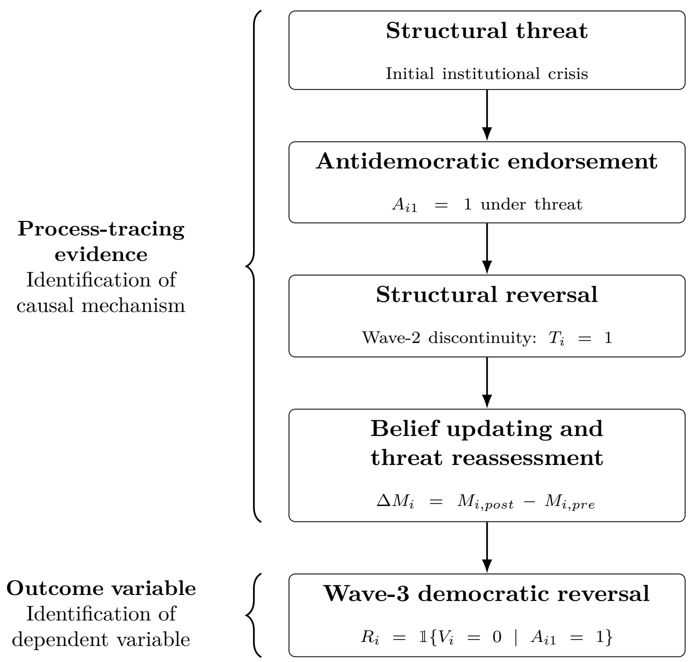

Figure 1: Process-tracing identification of the causal mechanism in the three-wave design.

Wave 1 establishes the initial state of the causal chain. Before exposure to campaign information, subjects report baseline preferences, attribute weights, perceived institutional capacity, perceived democratic-collapse risk, and standard democratic-support batteries. These measures serve three purposes. First, they parameterize the subject-specific utility function used by the incentive-compatible scoring rule in Wave 3 ("correct voting," see Lau and Redlawsk 2006). Second, they provide pre-discontinuity mechanism measures against which later updating can be assessed (see Figure 3). Third, they identify whether the subject endorses the antidemocratic option under threat.

define

After this baseline module, subjects enter a simulated polity characterized by institutional crisis. All subjects receive campaign information describing deteriorating structural conditions and candidate claims. The purpose of this common Wave-1 environment is to create the initial threat setting in which antidemocratic governance may become instrumentally attractive.

In Wave 1 *all* subjects receive an information stimuli describing an institutional crisis. The idea is to set the (fictitious) voting exercise (in Wave 3) within a context of institutional crisis. Accordingly, all subjects receive some information describing decaying structural conditions and candidate claims. Importantly, **information-acquiring is not imposed but rather endogenous, costly and incentive-compatible**: **subjects choose what kind of information buy to vote correctly** (if subjects vote correctly, they earn more points). Following the process-tracing logic of simulated campaign environments (Lau and Redlawsk 2006), subjects access information through a scrolling "news board" of labeled information boxes (issues, performance signals, institutional statements). **Clicking a box reveals content but incurs a small cost**; new boxes continue to appear, producing opportunity costs and forcing selective attention. Importantly, voters choose the kind of information they'd like to learn (self-selection). This mirrors the logic that campaign learning is shaped by constrained attention and endogenous information search (Lyons et al. 2016; Lau et al. 2013).

define election/candida profiles

define

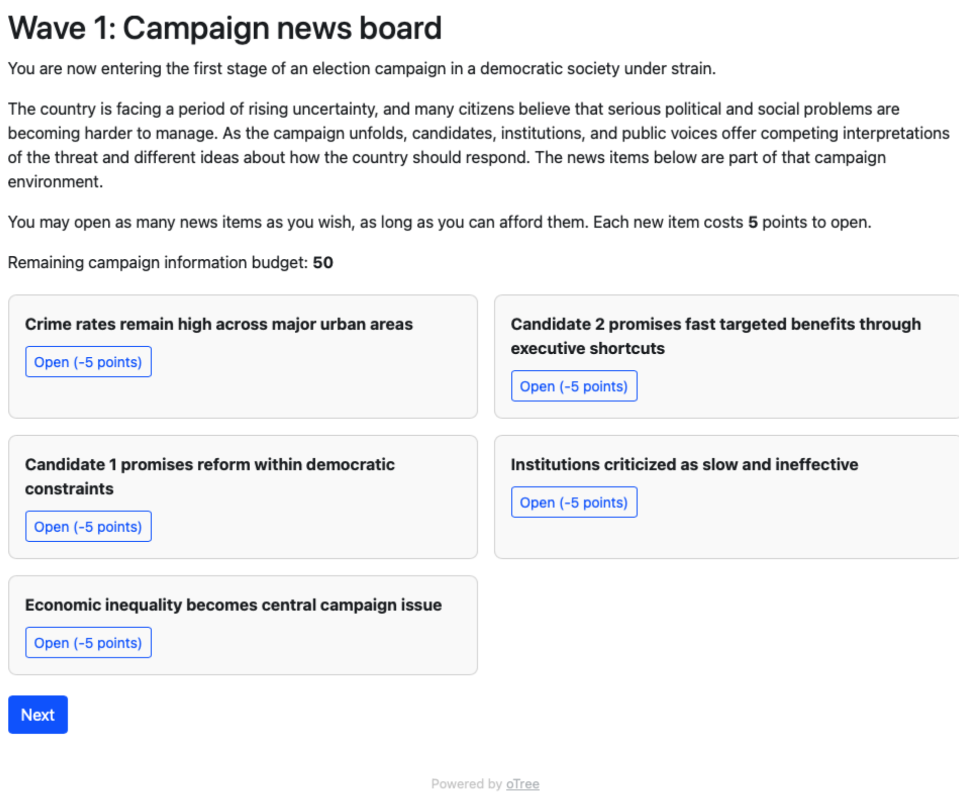

Figure 2: Example of News Board

Note: The figure illustrates the costly information-search environment used in Wave 1 of the campaign experiment. Participants enter a simulated election campaign under conditions of institutional and social strain and are given a limited campaign information budget. Each news item can be opened at a cost of 5 points, forcing participants to choose selectively which pieces of campaign information to acquire. The board includes information about structural threats, institutional performance, and candidate policy stances, allowing the design to trace how participants search for information before forming evaluations of democratic and anti-democratic alternatives.

crisis vignette

```
Wave 1: Institutional Crisis

The country is facing a prolonged institutional crisis.

Parliament has been unable to pass a national budget for eight months.
Several public services are operating under temporary funding rules.
Healthcare waiting times have doubled. Municipalities report uncertainty
about next year's services. Economic growth has stalled.

Two candidates are competing in the election. They offer different ways
of responding to the crisis, and new campaign information will become
available as the election unfolds.
```

Wave 2 implements an experimental discontinuity that enables me to devise a between-subject manipulation in the middle of the "campaign." Subjects are randomly assigned to (i) a **structural reversal treatment** in which structural conditions improve sharply (i.e., the threat presented in Wave 1 softens), or (ii) a **control** condition in which the adverse threat persists. The discontinuity is designed as an exogenous shock to the state of the world that **alters the theoretical necessity of democratic rule-breaking**. Critically, the candidates' *rhetoric* remains constant across conditions.

```
Wave 2: Institutional Crisis Softens

The institutional crisis has begun to ease.

After months of deadlock, parliament has reached a preliminary budget
agreement. Several public services are no longer operating under temporary
funding rules. Healthcare waiting times remain high, but they have stopped
increasing and new funding has been approved to reduce the backlog.
Municipalities report greater certainty about next year's services.
Economic growth remains weak, but recent indicators suggest that the
situation has stabilized.

The election campaign continues. The same two candidates remain in the race.
They still offer different ways of responding to the country's problems,
and new campaign information will become available as the election unfolds.
```

Wave 2 is therefore the mechanism-observation stage of the design. It does not provide the final outcome. Instead, it tests whether the structural discontinuity shifts the beliefs that, according to the theory, sustain antidemocratic tolerance under threat. If structural improvement reduces perceived democratic-collapse risk or increases perceived institutional capacity to deliver policy without norm violations, the design observes the intermediate movement expected by the causal mechanism before the Wave-3 vote occurs.

Wave 3 represents election day at the end of the hypothetical campaign and provides the primary substantive outcome. Subjects cast a final ballot between a pro-democratic candidate, who commits to institutional constraints and rule-of-law, and an antidemocratic candidate, who promises faster but norm-violating solutions to the structural problems framed in the experiment. The outcome

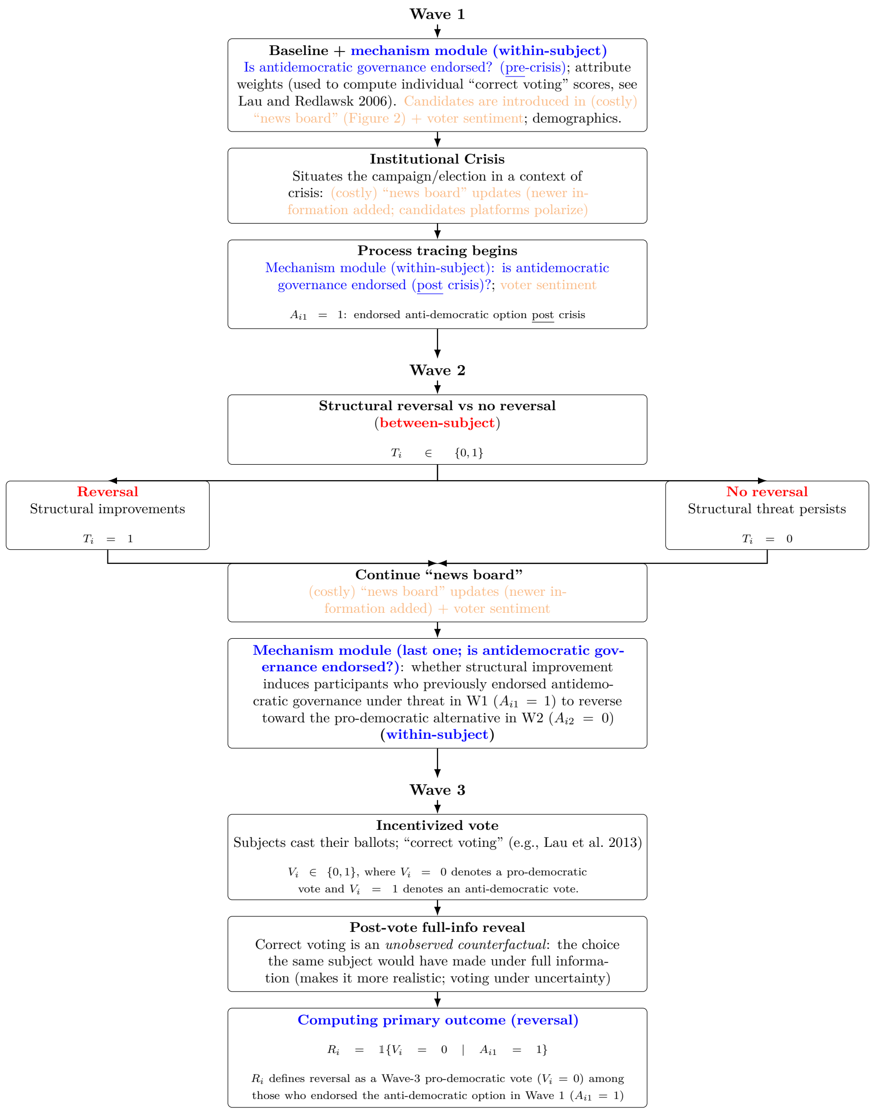

Figure 3: Three-wave panel voting lab experiment: within and between-subjects design.

variable is democratic reversal among subjects who endorsed the antidemocratic option under the Wave-1 threat environment. Thus, the key behavioral question is whether subjects who previously treated antidemocratic governance as instrumentally attractive ($A_{i1} = 1$) switch to the pro-democratic candidate once the structural conditions that justified such tolerance have changed ($R_i = \mathbb{1}\{V_i = 0 \mid A_{i1} = 1\}$).

"Correct voting" (Lau and Redlawsk 1997; Lau and Redlawsk 2006; Lau, Andersen, and Redlawsk 2008; Lau et al. 2013) is used instrumentally to make the task incentive compatible. Before campaign exposure in Wave 1, subjects state their preferred positions on different attributes and the importance weights attached to those attributes. These inputs define the subject's utility function. After the Wave-3 ballot, the experiment reveals the full information set and computes the candidate that maximizes the subject's own utility under full information.[1] Subjects receive an accuracy bonus if and only if their Wave-3 ballot matches that utility-maximizing benchmark. This scoring rule does not make correct voting a second substantive outcome. Its purpose is to make information search and voting consequential, reduce cheap talk, and ensure that the final ballot is cast under real incentives.

The design treats the laboratory as a controlled decision environment rather than as a physical location. This distinction matters because the experiment is designed to identify whether citizens revise antidemocratic tolerance when the structural conditions that made such tolerance instrumentally attractive are experimentally reversed. The relevant requirements are therefore not co-presence in a room, but control over assignment, information, timing, incentives, and payoff-relevant choice. This is precisely where online implementation is compatible with an economic laboratory logic. Online experiments can be "just as valid—both internally and externally—as laboratory and field experiments" when they preserve the conditions required for causal inference (Horton, Rand, and Zeckhauser 2011, p. 399); even interactive designs can produce data whose quality is "adequate and reliable" when attrition, comprehension, and payment rules are properly managed (Arechar, Gächter, and Molleman 2018, p. 100). More recent evidence comparing otherwise identical laboratory and online implementations reaches the same conclusion: Prissé and Jorrat (2022, p. 8) argue that "lack of control is not an issue" for several standard economic tasks. The proposed design therefore uses an online laboratory format to study a political-psychological process–learning, memory, and threat reassessment–under economic-experimental conditions: costly information acquisition, incentive-compatible voting, and a candidate-neutral payoff rule tied to subject-specific utility.

A note on implementation The memory literature implies that the effect of Wave-1 threat information should decay unless it is refreshed or remains highly accessible, and that what subjects remember at Wave 3 may depend on whether they are engaging in memory-based evaluation or a more durable online tally process (Hill et al. 2013). Likewise, framing effects can persist across delayed time points but vary systematically by political knowledge, underscoring why repeated measurement across waves is necessary rather than reliance on immediate post-treatment responses (Lecheler and De Vreese 2011). Wave 2 therefore introduces new campaign information in the board (see Figure 2) and rephrases old information (maintaining costs). Importantly, it keeps the same

[1] I reveal full information *after* voting to address an identification problem: correct voting is defined relative to an *unobserved counterfactual* (the choice the same subject would have made under full information; Lau and Redlawsk 2006). The design therefore must create a full-information benchmark that is observed only *after* the ballot is cast, so the benchmark can be used to classify the earlier, uncertainty-based vote rather than determine it. If full information were provided pre-vote, the design would collapse to a single fully informed choice and the "correctness" counterfactual would no longer be separately identifiable. Second, for incentives, withholding complete information preserves a realistic campaign decision problem in which information acquisition is costly and uncertainty remains, preventing the task from collapsing into cheap talk or a trivial maximization with guaranteed full knowledge before voting.

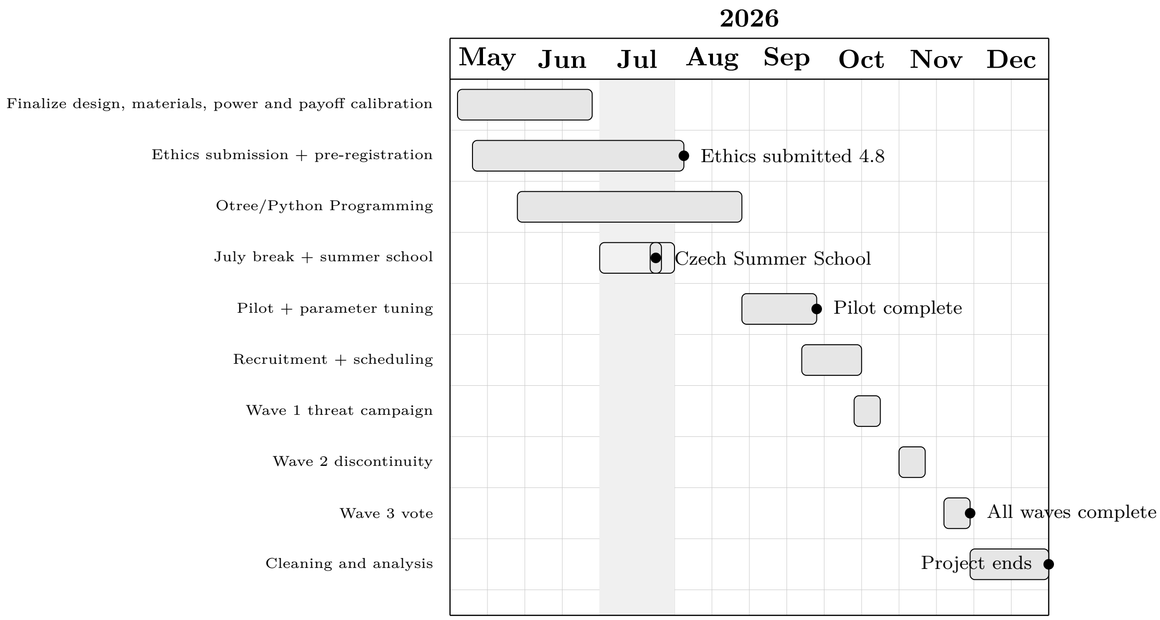

Figure 4: Timeline, May—December 2026. The schedule assumes ethical review submission by August 4 and project completion by the end of December 2026.

module on democratic attitudes asked in Wave 1. This also allows us to test whether subjects revise or instead persist in an incorrect status despite feedback from the changed environment (Esponda, Vespa, and Yuksel 2024).

Because the focus is *learning and reversal in a campaign over time*, the waves must be separated enough to allow memory decay and differential accessibility to operate, rather than collapsing everything into an immediate post-treatment window. Accordingly, we run Wave 1, Wave 2, and Wave 3 as three short lab sessions separated by 7 days. A 1-week delay is the critical threshold where competitive framing effects become statistically significant and interpretable. (Lecheler and Vreese 2013, p. 158) note that shorter delays (like 1 day) suffer from recency effects. Thus, a 1-week interval provides optimal conditions for observing the full dynamics of attitude change (Lecheler and De Vreese 2011) (i.e., not too recent, not too long ago).

Spacing also increases the inferential leverage of process tracing: if reversals occur only when Wave-2 improvement is salient and remembered at Wave 3, that pattern is consistent with memory-based evaluation; if reversals persist despite delay, that is consistent with more durable updating (an online tally or revised mental model) (Hill et al. 2013; Esponda, Vespa, and Yuksel 2024; Levy and Razin 2025). Substantively, this distinction matters for democratic resilience: **if commitment to democratic norms depends on what remains cognitively accessible from recent events, then improvements may generate only fragile, short-lived "resilience**," whereas persistence after delay suggests that citizens have internalized institutional constraints as part of their evaluative standard rather than merely reacting to transient campaign salience.

Quantities of Interest

The design distinguishes among three components: the primary substantive outcome, the instrumental scoring rule, and the process-tracing observations. The primary substantive outcome is Wave-3 democratic reversal among subjects who endorsed the antidemocratic option under the Wave-1 threat environment. Correct voting is not a second substantive outcome; it is an incentive-compatible scoring rule used to make information search and voting consequential. Wave-1 and Wave-2 measures are mechanism-stage observations used to trace the causal sequence leading to Wave-3 reversal.

Define subject i's final vote in Wave 3 as $V_i \in \{0,1\}$, where $V_i = 1$ denotes voting for the antidemocratic candidate and $V_i = 0$ denotes voting for the pro-democratic candidate. Let $T_i \in \{0,1\}$ indicate assignment to the Wave-2 structural discontinuity, with $T_i = 1$ for the structural reversal condition and $T_i = 0$ for the no-reversal control. Let $A_{i1} \in \{0,1\}$ denote whether subject i endorsed the antidemocratic option in Wave 1 after exposure to the common threat environment. This quantity is not the final outcome. It identifies the initial state of the causal sequence and defines the subgroup for whom democratic reversal is theoretically meaningful.

Primary outcome: Wave-3 democratic reversal. For subjects with $A_{i1} = 1$, define

$$R_i \equiv \mathbb{1}\{V_i = 0\}.$$

Thus, $R_i = 1$ indicates that a subject who endorsed the antidemocratic option under the Wave-1 threat environment casts a pro-democratic vote in Wave 3, while $R_i = 0$ indicates persistence. Let $R_i(1)$ and $R_i(0)$ denote potential reversal outcomes under assignment to the structural reversal condition and the no-reversal control, respectively. The primary causal estimand is

$$CATE_R \equiv E\left[R_i(1) - R_i(0) \mid A_{i1} = 1\right].$$

This estimand captures whether structural improvement induces citizens who previously endorsed antidemocratic governance under threat to reverse toward the pro-democratic alternative at the final election.

Instrumental scoring rule: correct voting. Correct voting is used to incentivize consequential information search and voting. Let $c_i^\star$ denote the candidate that maximizes subject i's own utility under full information:

$$c_i^\star = \arg\max_c U_i(c),$$

where $U_i(c)$ is parameterized by the preferences and attribute weights elicited before campaign exposure in Wave 1. After the Wave-3 ballot, the experiment reveals the full information set and computes $c_i^\star$. Subjects receive an accuracy bonus when their observed Wave-3 vote matches this benchmark:

$$C_i = \mathbb{1}\{V_i = c_i^\star\}.$$

This quantity is used for payment and incentive compatibility. It is not the main dependent variable. Its function is to make information acquisition costly but potentially valuable, reduce cheap talk, and ensure that the final ballot is cast under a consequential decision environment.

Process-tracing observations. The mechanism measures are elicited before and after the Wave-2 discontinuity to evaluate whether the hypothesized causal chain is empirically visible. Let $M_{i,pre}$ denote a mechanism measure observed before the discontinuity and $M_{i,post}$ the same measure observed after the discontinuity. Key mechanism measures include perceived institutional capacity to deliver policy without norm violations and perceived risk of democratic collapse. Updating is defined as

$$\Delta M_i = M_{i,post} - M_{i,pre}.$$

The process-tracing analysis asks whether T_i shifts ΔM_i in the theoretically expected direction and whether these changes are associated with the primary Wave-3 outcome R_i. This evidence does not replace the outcome. It evaluates whether the treatment effect on final behavior is accompanied by the anticipated intermediate steps: structural improvement, belief updating, threat reassessment, and final democratic reversal.

4 Statistical considerations and power

Across published candidate-choice and conjoint experiments, norm violations and related cues shift support by anything from low single digits to the mid-teens (and, in some settings, substantially more). In Canada/US intra-party choice tasks, the penalty for an undemocratic candidate attribute can be as small as 1.6 points (ignoring unfavorable court decisions) (Gidengil, Stolle, and Bergeron-Boutin 2021, p. 20). In five-country conjoint evidence, undemocratic behavior reduces support by -0.06 to -0.20 on a five-point support scale (Frederiksen 2022, p. 1149). In Finland, the design is powered to detect AMCE ≈ 0.03, i.e., about a 3 percentage-point change in leader favorability (Saikkonen and Christensen 2023, p. 6). In the US donor/voter conjoint, co-partisan support drops from 84% (no violations) to 62% (three violations), with one violation reducing support to 74–80% and donors exhibiting net decreases of 30–45 percentage points when moving from zero to three violations (Carey et al. 2022, p. 237). In Hungary, certain leader-profile attributes reduce selection probabilities by 7%–7.5% (openness to new influences), while subgroup differences in tolerance for two democratic attributes are about 2% in each case (Wunsch and Gessler 2023, p. 924). In party-conditioned models, being in power increases support for norm-eroding policies by 2 percentage points and perceived out-group threat increases support by 4 percentage points (Simonovits, McCoy, and Littvay 2022, p. 1809). In candidate-choice evidence on polarization, undemocratic reforms carry an estimated 16.07% electoral penalty and "candidates supporting undemocratic reforms" obtain about 5% fewer votes (Svolik 2020, pp. 164–166). Finally, when an illiberal candidate wins, support for an executive-led coup is 62% among winners versus 31% among losers (a 31 percentage-point gap) (Cohen et al. 2023, p. 261). **Taken together, this heterogeneous evidence implies that both baseline endorsement rates and treatment effects can vary widely across institutional contexts, partisan environments, and the severity or accumulation of violations—which is precisely the scenario this design is built to study**.

Anchoring power calculations in this literature, I use a planning value of $p(A_{i1} = 1) \approx 0.40$ for Wave-1 anti-democratic endorsement (i.e., a 60/40 split between non-endorsers and endorsers, with a plausible range of 0.30–0.50; see Figure 5). The primary outcome is binary: Wave-3 democratic reversal (R_i) among Wave-1 antidemocratic endorsers ($A_{i1} = 1$). Correct voting (C_i) is not treated as a substantive outcome for power purposes. It is used as an incentive-compatible scoring rule that makes information search and voting consequential.

The main estimands will be estimated with logistic regression. With a two-sided $\alpha = 0.05$ and $1 - \beta = 0.80$, I will recruit $N \approx 350$ participants to obtain an expected final analytic sample of $N \approx 300$ after approximately 50 dropouts. With $p(A_{i1} = 1) \approx 0.40$, the effective sample for the primary outcome is approximately $n \approx 60$ Wave-1 antidemocratic endorsers per arm. This design is therefore powered to detect large reversal effects, on the order of roughly 25 percentage points. **The resulting risk is that if the true effect in this design is closer to the mid-teen range suggested by several studies above, the experiment may still be underpowered for conventional hypothesis testing**, and if the realized endorsement rate is closer to the lower end of the plausible range (e.g., 0.30), power for R_i will fall further.

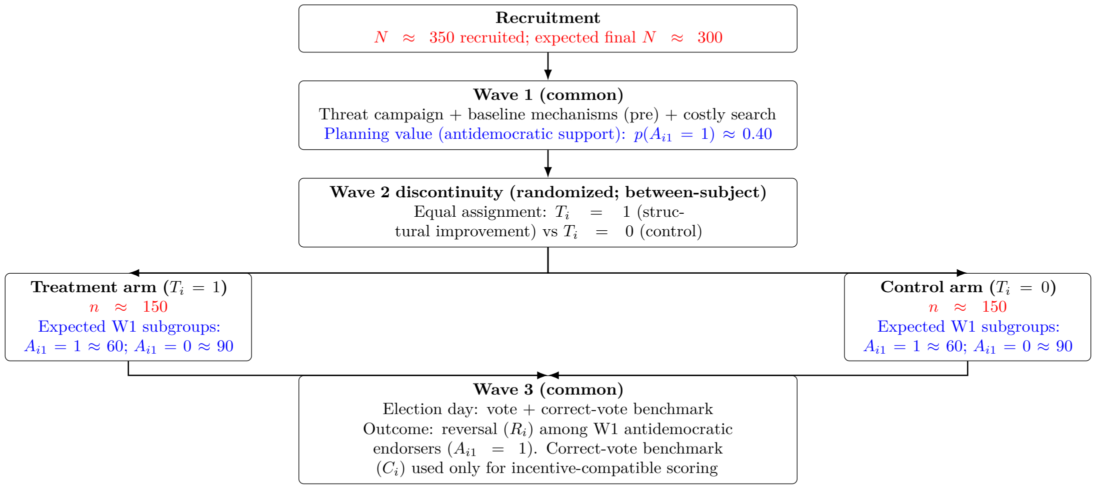

Figure 5: Planned sample allocation ($N \approx 350$ recruited; expected final $N \approx 300$ after approximately 50 dropouts). Within each randomized arm, the expected Wave-1 subgroup counts are based on the planning value $p(A_{i1} = 1) \approx 0.40$: approximately 60 Wave-1 anti-democratic endorsers and 90 non-endorsers per arm.

5 Ethical Considerations and Mitigation

The design intentionally exposes participants to environments that make antidemocratic tolerance appear instrumentally reasonable before reversing those conditions. This creates ethical obligations, but it also makes the design scientifically valuable. Participants are informed that the study focuses on political attitudes and learning. The post-experimental debrief clarifies that shifts in participants' judgments are interpreted not as inconsistency, but as evidence of adaptive democratic learning.

References

Ananda, Aurelia and Damien Bol (Aug. 2021). "Does Knowing Democracy Affect Answers to Democratic Support Questions? A Survey Experiment in Indonesia." In: *International Journal of Public Opinion Research* 33.2, pp. 433–443.

Arechar, Antonio A., Simon Gächter, and Lucas Molleman (Mar. 2018). "Conducting interactive experiments online." In: *Experimental Economics* 21.1, pp. 99–131.

Bennett, Andrew and Jeffrey Checkel (Nov. 2014). "Process Tracing: From philosophical roots to best practices." In: *Process Tracing*. Cambridge University Press, pp. 3–38.

Capoccia, Giovanni (2007). *Defending Democracy*. Johns Hopkins University Press.

Carey, John et al. (2022). "Who Will Defend Democracy? Evaluating Tradeoffs in Candidate Support Among Partisan Donors and Voters." In: *Journal of Elections, Public Opinion and Parties* 32.1, pp. 230–245.

Cohen, Mollie et al. (2023). "Winners' Consent? Citizen Commitment to Democracy When Illiberal Candidates Win Elections." In: *American Journal of Political Science* 67.2, pp. 261–276.

Collier, David (Oct. 2011). "Understanding Process Tracing." In: *PS: Political Science & Politics* 44.4, pp. 823–830.

Esponda, Ignacio, Emanuel Vespa, and Sevgi Yuksel (Mar. 2024). "Mental Models and Learning: The Case of Base-Rate Neglect." In: *American Economic Review* 114.3, pp. 752–782.

Frederiksen, Kristian Vrede Skaaning (2022). “Does Competence Make Citizens Tolerate Undemocratic Behavior?” In: *American Political Science Review* 116.3, pp. 1147–1153.
Gerring, John (June 2004). “What Is a Case Study and What Is It Good for?” In: *American Political Science Review* 98.2, pp. 341–354.
Gidengil, Elisabeth, Dietlind Stolle, and Olivier Bergeron-Boutin (2021). “The Partisan Nature of Support for Democratic Backsliding: A Comparative Perspective.” In: *European Journal of Political Research*.
Graham, Matthew and Milan Svolik (2020). “Democracy in America? Partisanship, Polarization, and the Robustness of Support for Democracy in the United States.” In: *American Political Science Review* 114.2, pp. 392–409.
Grillo, Edoardo and Carlo Prato (2023). “Reference Points and Democratic Backsliding.” In: *American Journal of Political Science* 67.1, pp. 71–88.
Hill, Seth et al. (Oct. 2013). “How Quickly We Forget: The Duration of Persuasion Effects From Mass Communication.” In: *Political Communication* 30.4, pp. 521–547.
Horton, John J., David G. Rand, and Richard J. Zeckhauser (Sept. 2011). “The online laboratory: conducting experiments in a real labor market.” In: *Experimental Economics* 14.3, pp. 399–425.
Kingzette, Jonathan et al. (2021). “How Affective Polarization Undermines Support for Democratic Norms.” In: *Public Opinion Quarterly* 85.2, pp. 663–677.
Lau, Richard and David Redlawsk (Sept. 1997). “Voting Correctly.” In: *American Political Science Review* 91.3, pp. 585–598.
Lau, Richard and David Redlawsk (2006). *How Voters Decide, Information Processing in Election Campaigns*. Cambridge University Press.
Lau, Richard et al. (Mar. 2013). “Correct Voting Across Thirty-Three Democracies: A Preliminary Analysis.” In: *British Journal of Political Science* 44.02, pp. 239–259.
Lau, Richard R., David Andersen, and David Redlawsk (Apr. 2008). “An Exploration of Correct Voting in Recent U.S. Presidential Elections.” In: *American Journal of Political Science* 52.2, pp. 395–411.
Lecheler, Sophie and Claes De Vreese (2011). “Getting Real: The Duration of Framing Effects.” In: *Journal of Communication* 61.5, pp. 959–983.
Lecheler, Sophie and Claes de Vreese (Apr. 2013). “What a Difference a Day Makes? The Effects of Repetitive and Competitive News Framing Over Time.” In: *Communication Research* 40.2, pp. 147–175.
Levitsky, Steven and Daniel Ziblatt (2018). *How Democracies Die*. Crown.
Levy, Gilat and Ronny Razin (Feb. 2025). “Political Social Learning: Short-Term Memory and Cycles of Polarization.” In: *American Economic Review* 115.2, pp. 635–659.
Levy, Jack (Feb. 2008). “Case Studies: Types, Designs, and Logics of Inference.” In: *Conflict Management and Peace Science* 25.1, pp. 1–18. arXiv: `arXiv:1011.1669v3`.
Lyons, Jeffrey et al. (July 2016). “Personality, Interpersonal Disagreement, and Electoral Information.” In: *The Journal of Politics* 78.3, pp. 806–821.
Mahoney, James (Dec. 2010). “After KKV: The New Methodology of Qualitative Research.” In: *World Politics* 62.1, pp. 120–147.
Morton, Rebecca and Kenneth Williams (2010). *Experimental Political Science and the Study of Causality: From Nature to the Lab*. Cambridge University Press.
Neundorf, Anja et al. (Dec. 2025). “Elections Without Constraints? The Appeal of Electoral Autocracy Across the World.” In: *British Journal of Political Science* 56.
Orhan, Yunus Emre (2022). “The Relationship between Affective Polarization and Democratic Backsliding: Comparative Evidence.” In: *Democratization* 29.4, pp. 714–735.

Pearl, Judea (Feb. 2015). “Causes of Effects and Effects of Causes.” In: *Sociological Methods & Research* 44.1, pp. 149–164.

Prissé, Benjamin and Diego Jorrat (Oct. 2022). “Lab vs online experiments: No differences.” In: *Journal of Behavioral and Experimental Economics* 100, p. 101910.

Redlawsk, David (Feb. 2001). “You Must Remember This: A Test of the On-Line Model of Voting.” In: *The Journal of Politics* 63.1, pp. 29–58.

Saikkonen, Inga A.-L. and Henrik Serup Christensen (Mar. 2023). “Guardians of Democracy or Passive Bystanders? A Conjoint Experiment on Elite Transgressions of Democratic Norms.” In: *Political Research Quarterly* 76.1, pp. 127–142.

Simonovits, Gabor, Jennifer McCoy, and Levente Littvay (2022). “Democratic Hypocrisy and Out-Group Threat: Explaining Citizen Support for Democratic Erosion.” In: *Journal of Politics* 84.3, pp. 1806–1811.

Svolik, Milan (2020). “When Polarization Trumps Civic Virtue: Partisan Conflict and the Subversion of Democracy by Incumbents.” In: *Quarterly Journal of Political Science* 15.1, pp. 3–31.

Voelkel, Jan G. et al. (2024). “Megastudy testing 25 treatments to reduce antidemocratic attitudes and partisan animosity.” In: *Science* 386.6719, eadh4764.

Wunsch, Natasha and Theresa Gessler (July 2023). “Who Tolerates Democratic Backsliding? A Mosaic Approach to Voters’ Responses to Authoritarian Leadership in Hungary.” In: *Democratization* 30.5, pp. 914–937.

.... **Word count**: 4,194 ....