

Undemocratic Reversal: Crisis, Recovery, and Weakly Constrained Executive Power in a Laboratory Public-Goods Game

HECTOR BAHAMONDE*^{1,2}

¹INVEST Research Flagship Centre, University of Turku, Finland

²Social Science Research Institute, Åbo Akademi University, Finland

Working proposal — July 23, 2026

Abstract

PENDING

See [prototype here](https://panel-lab-experiment.onrender.com): <https://panel-lab-experiment.onrender.com>

*hector.bahamonde@utu.fi; hectorbahamonde.com

1 Research Question and Contribution

Democratic backsliding research has established that citizens can learn to tolerate norm violations when structural threats make such tolerance instrumentally attractive. The focus of this proposal is whether that learning is reversible when those structural conditions improve. Thus, the most intriguing finding in recent democratic backsliding research is not that citizens tolerate norm violations (Graham and Svolik 2020). My claim is that they *learn* to do so. As citizens secularly experience norm-violating leaders and accumulate information about non-democratic behaviors (Levitsky and Ziblatt 2018), they recalibrate expectations, adjust their definitions of democracy itself (Ananda and Bol 2021), and gradually become “normalized” to authoritarianism (Grillo and Prato 2023). Yet this learning process raises a critical question: **Can that undemocratic learning be reversed?** If citizens can learn authoritarian tolerance, can they *un-learn* it when the institutional conditions and threats that gave rise to such undemocratic leaders change? And critically, which institutional conditions trigger such reversals? These questions are fundamental for democratic resilience, and importantly, for *defending* democracy (Capoccia 2007).

Thus, this study focuses on this behavioral transition in the context of *undemocratic reversal*. It is not simply low support for an undemocratic alternative, nor a between-person difference in democratic attitudes. It is an observed within-person change: a participant initially endorses an institution that concentrates authority and weakens ex ante control, then later chooses an institution that restores collective choice and constraints.

The design contributes to the literature in three ways. First, it turns democratic constraint into a set of payoff-relevant rules governing interaction. Participants do not merely rate political descriptions (vignettes). Their ballots select who controls endowments, how a public good is produced, and whether an executive may extract a private rent. This embeds democratic trade-offs in the public-good provision, delegation and voting literatures (Hamman, Weber, and Woon 2011).

Second, it identifies a structural cause of reversal. All participants encounter a public-service crisis in Block 1: **“Imagine that you are a citizen in a society where public services have been under serious strain for a long time. Parliament has failed to pass a budget for eight months, healthcare waiting times have doubled, and planning for public services has stalled.”** For Block 2, matched pools are randomly assigned either to recovery, in which the crisis recedes, or to persistence, in which it continues. The leader-led institution carries the same multiplier, authority, and personal-transfer rule across both blocks. The treatment therefore isolates the crisis-created productivity premium attached to concentrating power.

Third, it identifies changes in *revealed institutional preference*. Democratic commitments acquire their meaning through consequential choices. During the crisis, a vote for one-person rule can express a willingness to exchange collective control for productive performance. Recovery changes that bargain by raising citizen-controlled productivity to the leader-led rate. Leader control still offers compulsory coordination: one person can require equal allocations and thereby overcome voluntary free-riding. A citizen who returns to collective choice therefore reveals a preference for democratic control even when one-person rule may coordinate provision more successfully. This is a demanding form of democratic reversal: the same citizen accepts the strategic risks of collective choice once the crisis-created productivity premium for concentrating power has expired.

In simple terms

This proposal asks whether citizens remain committed to democratic control when a crisis makes one-person rule appear faster and more effective. The undemocratic choice occurs when citizens knowingly accept a rule under which one person replaces collective choice and can impose a decision on everyone without prior group approval. The option to move points into a personal account gives concentrated authority a payoff-relevant temptation and makes its consequences material. The game thereby reveals how citizens and selected leaders behave when collective control competes with effectiveness and individual reward.

The experiment then asks whether support for one-person rule survives after the crisis recedes. Recovery directly changes the payoff structure: the return from citizen-controlled provision rises from 1.50 to 2.50 and matches the leader-led rate. The leader retains the ability to require equal allocations and coordinate provision, together with the option to move fund points into a personal account. Citizens who return to collective decision making therefore reclaim democratic control while accepting the continuing risk of voluntary free-riding. In short: *when one-person rule loses its crisis-created productivity advantage, do citizens value collective control enough to take it back?*

2 Theory

2.1 Conditional tolerance of unconstrained authority

Citizens need not reject democratic constraints in principle to support an unconstrained executive in practice. They may instead accept discretion conditionally, i.e., when institutions appear incapable of producing valued outcomes. This logic is compatible with evidence that electoral and material considerations can outweigh checks on executives and that citizens interpret democratic violations through partisan and performance-based trade-offs (Kingzette et al. 2021; Saikkonen and Christensen 2023; Wunsch, Jacob, and Derksen 2025). It also clarifies why democratic tolerance can be learned: repeated evidence of institutional failure changes the expected payoff from retaining constraints.

In the experiment, the citizen-led procedure (D) preserves individual control: each person decides how many points to place in a public-services fund, and no one can remove points from that fund for personal use. Its low Block-1 productivity, however, makes public-service provision difficult. The leader-led procedure (A) represents a faster coordinated response: one person makes the decision for all five citizens, and the remaining fund is converted into public services more effectively. The same leader may also move part of the fund to a personal account. The productive multiplier gives this faster-response advantage a consequential payoff counterpart. Support for A can therefore express an instrumental response to poor performance during a prolonged public-service crisis.

H1: Crisis endorsement. In Block 1, support for the leader-led method will increase when participants' experience of its productivity premium and compulsory coordination outweighs the risk of a personal transfer.

2.2 Authorization and enactment of backsliding

Democratic backsliding involves more than the presence of a self-interested leader. It also depends on whether citizens authorize arrangements that weaken collective control and on what officeholders do with the discretion they receive. The experiment separates these behavioral margins. A ballot for A authorizes one participant to decide for everyone, weakens ex ante control over the group's

resources, and accepts the possibility of personal appropriation. Conditional on being selected as leader, the transfer ρ records whether and how extensively that participant uses the granted discretion for personal gain.

This distinction connects the public-good game to democratic backsliding without treating a stylized allocation as literal corruption. Voting for A is an *enabling* act; moving public-service points to a personal account is an *enactment* of the distributive risk that weakened constraints create. Subsequent ballots then show how participants respond to both institutional performance and observed misuse. The design can therefore study individual adherence to collective control in a transparent incentive environment, rather than asking respondents whether appropriation is likely in contemporary Finland.

H2: Use of discretion. Some selected leaders will move public-service points to their personal accounts, and larger observed transfers will be followed by lower support for the leader-led method among exposed participants.

2.3 Structural updating and democratic reversal

The central theoretical claim is conditional. If support for leader discretion was learned from weak structural performance, then credible improvement should change its relative appeal. Before Block 2, recovery raises the citizen-led method’s productivity to the leader-led method’s level. This removes the exogenously imposed productivity premium for one-person rule while preserving its compulsory coordination and the leader’s discretion to move fund points into a personal account. Citizens can then transform every contributed point into public services at the same rate under either method, while only the leader can require everyone to contribute. The resulting choice retains real economic stakes: collective control carries the risk of free-riding, whereas leader control carries the loss of individual choice and the possibility of a personal transfer. Returning to D under these conditions reveals a willingness to restore democratic control despite the leader-led method’s remaining coordination advantage.

The persistence condition supplies the counterfactual. All participants remain in the same session, receive a Block-2 description, and play the same number of additional rounds; only the productivity of the citizen-led method differs. The treatment contrast therefore separates recovery-induced reversal from additional experience, fatigue, end-game behavior, and regression to the mean. The design identifies structural updating and within-session reversal with no return-visit attrition.

H3: Structural updating. Participants assigned to recovery will report a larger increase in perceived citizen-led effectiveness and a larger reduction in the perceived need for leader control than participants assigned to persistence.

H4: Democratic reversal. Among participants who choose A in Block 1’s final round, assignment to recovery will increase the probability of choosing D in Block 2’s final round despite the leader-led method’s remaining coordination advantage.

H5: Behavioral path. Recovery will change the path of institutional choice across Block 2: support for D should rise once its restored productivity is disclosed, while exposure to leader appropriation should make continued delegation less attractive.

3 Experimental Design

3.1 Overview

The study consists of one synchronized laboratory session lasting approximately 75–90 minutes. Participants complete two consecutive ten-round blocks and remain in stable ten-person matching pools throughout. Before each paid round, the software anonymously divides every pool into two temporary five-person groups. The Block-2 condition is randomized at the pool level and disclosed between blocks. Participants therefore interact repeatedly within a treatment-consistent local population but cannot build a reliable bilateral reputation.




Figure 1: One-session experimental sequence. Every paid round contains a secret institutional ballot, an institution-contingent allocation decision, feedback, and anonymous rematching.




Figure 2: Matching pools, anonymous rematching, and Block-2 treatment assignment.

In simple terms

The design uses permanent ten-person matching pools, denoted by g . Before each paid round, the software divides every pool into two temporary five-person groups that vote and play separately. These groups dissolve after the round and are formed again before the next one. Participants never interact outside their pool. Accordingly, the colors in Figure 2 show current-round group membership, while p_i identifies the same participant throughout the session. Before Block 2, the entire pool learns that the public-service crisis has either receded ($T_g = 1$) or persists ($T_g = 0$). The planned sample repeats this structure across 40 pools.

This arrangement creates interaction without fixed teams. Its principal cost is repeated exposure within pools, but anonymity and rematching limit partner-specific reputation, collusion, and retaliation while preserving social learning. Stable pools also keep potential spillovers within treatment-homogeneous environments, provide 40 independently randomized clusters, and permit two complete five-person groups in every round.

Statistically, democratic reversal (R_i) is measured for each participant p_i , whereas treatment (T_g) is assigned to the pool. Individual outcomes are therefore compared across treatment conditions with standard errors clustered by pool. The temporary five-person groups organize play but are not independent treatment units.

3.2 The institutional-choice public-good game

Each round begins with 20 points for each of five group members. Every member privately votes between the citizen-led procedure (D) and the leader-led procedure (A). A simple majority of at least three votes selects the binding method. Individual ballots remain secret; the group sees the aggregate vote and the selected rule. The neutral labels shown to participants are “Citizens decide” and “A leader decides.”

Citizen-led procedure (D). Each member chooses an allocation c_i of 0–20 points to a public-services fund; unallocated points remain in the member’s personal account. Let m_{gb} denote the fund’s productivity for matching pool g in block b . Player i earns

$$\pi_{igbr}^D = 20 - c_{igbr} + \frac{m_{gb}}{5} \sum_{j=1}^5 c_{jgbr}. \quad (1)$$

The marginal private return is below 1 in both public-service environments, preserving the social dilemma, while the group return exceeds 1. In Block 1, $m_{g1} = 1.50$ for everyone. In Block 2, $m_{g2} = 2.50$ under recovery and $m_{g2} = 1.50$ under persistence.

Leader-led procedure (A). One group member is selected as leader. Selection is balanced toward members who have served least often, with random resolution of ties. The leader chooses the number of points $t \in \{0, \dots, 20\}$ that every citizen must place in the public-services fund. The leader may then move $\rho \in \{0, \dots, 20\}$ fund points to a personal account, subject to $\rho \leq 5t$. The remaining fund $5t - \rho$ always uses multiplier 2.50. A nonleader earns

$$\pi_{igbr}^A = 20 - t + \frac{2.50}{5}(5t - \rho), \quad (2)$$

while the leader earns

$$\pi_{igbr}^{A,L} = 20 - t + \frac{2.50}{5}(5t - \rho) + \rho. \quad (3)$$

All members observe the required allocation, the amount moved to the leader’s personal account, the public-service return, and their own payoff. Leader control can coordinate provision and, in Block 1, produces public services more effectively. It also creates a distributive conflict grounded in another participant’s decision.

How incentives map onto reversal. The treatment implements the changing social context through the relative productivity of the two decision methods. During the crisis, the citizen-led fund uses $m = 1.50$, whereas the leader-led fund uses $m = 2.50$. Under recovery, the citizen-led multiplier rises to 2.50 and the leader-led rule remains unchanged. The persistence condition retains the crisis incentive structure in Block 2, providing the counterfactual for changes caused by repeated play, fatigue, or end-game behavior.

Table 1 illustrates the comparison by supposing that ten of each citizen’s twenty points enter the fund and that the leader makes no personal transfer. Holding the fund allocation fixed isolates the multiplier, with ten points serving as a common reference allocation.

Table 1: Illustrative payoffs across decision methods and crisis conditions

Decision method	Crisis: leader productivity premium	Recovery: equal productivity
Citizens decide (D)	$m = 1.50$. If every citizen places 10 points in the fund, each earns $\pi_i^D = 20 - 10 + \frac{1.50}{5}(5 \times 10) = 25.$	$m = 2.50$. With the same choices, each earns $\pi_i^D = 20 - 10 + \frac{2.50}{5}(5 \times 10) = 35.$
A leader decides (A)	$m = 2.50$. If the leader requires 10 points and sets $\rho = 0$, each earns $\pi_i^A = 20 - 10 + \frac{2.50}{5}(5 \times 10) = 35.$	$m = 2.50$, unchanged. The same leader decision again gives each citizen $\pi_i^A = 35.$

Note: The persistence condition retains the crisis column in Block 2. If the leader moves points to a personal account ($\rho > 0$), nonleaders receive less and the leader receives that transfer in addition to the common return, as defined in Equations (2) and (3).

The matrix maps the incentive structure onto the hypotheses. In the example, the leader-led method produces a ten-point payoff advantage during the crisis, and recovery equalizes the productivity rate at $m = 2.50$. At that common rate, a point voluntarily placed in the citizen-led fund returns $2.50/5 = 0.50$ points to the citizen who placed it. The private incentive to free-ride therefore remains, while a leader can coordinate provision by requiring equal allocations. Leader control combines this coordination advantage with authority over individual choices and access to a personal transfer. For example, if the leader requires ten points and moves ten fund points to a personal account, each nonleader earns 30, the leader earns 40, and a citizen’s ex ante expected payoff before leader selection is 32 under symmetric selection probabilities.

Accordingly, H1 concerns support for leader control when it combines a productivity premium with compulsory coordination. H4 concerns whether citizens return to collective control once that experimentally assigned premium expires and the leader’s coordination advantage remains.

Because both methods use 2.50 under recovery, a recovery-induced increase in D ballots records a consequential preference for collective control under continuing free-riding risk. The persistence pools identify this change in revealed institutional preference against an environment in which leader control retains both advantages.

In simple terms

Each round starts with 20 points for each citizen. The five citizens first cast private ballots ($V_{ibr} \in \{D, A\}$) about who should make the public-service decision. If at least three vote for “Citizens decide” (D), each citizen chooses how many points to put in the public-services fund (c_{igbr}). The software adds those points, multiplies the fund by the current rate (m_{gb}), and divides the result equally among all five citizens. Points not put in the fund remain in each citizen’s personal account.

If at least three vote for “A leader decides” (A), the software appoints one group member, randomly selecting among those who have served least often. Participants therefore choose the decision method, not a particular candidate; there is no campaign. The leader decides how many points every citizen must put in the fund (t). The leader may then move up to 20 of those fund points to a personal account ($\rho \leq \min\{20, 5t\}$). The points left in the fund ($5t - \rho$) are multiplied by 2.50 and divided equally among all five citizens; the leader also keeps the amount moved.

At the end of the round, everyone sees the selected method, the fund decisions, any amount moved by the leader, the public-service return, and their own payoff (π_{igbr}^D under D or π_{igbr}^A under A). The temporary group then dissolves, and participants are anonymously rematched within their ten-person pool (g).

Table 2: Implemented design parameters

Feature	Implemented value
Laboratory session	One synchronized visit, approximately 75–90 minutes
Paid rounds	10 per block (20 total), preceded by unpaid practice. (pre-reg 1st round; “hot decision.”)
Group and matching pool	5-person groups; 10-person treatment-matched pools maybe it builds reputation? think; find reference
Rematching	Anonymous rematching within pool before every paid round
Round endowment	20 points per participant
Citizen-led multiplier	1.50 in Block 1; 2.50 recovery / 1.50 persistence in Block 2
Leader-led multiplier	2.50 in both blocks and conditions
Leader’s personal transfer	0–20 fund points, limited by the number of points in the fund
Performance payment	One randomly selected game round from each block (find reference)
Exchange rate	EUR 0.20 per point
Show-up payment	EUR 10.00

The random-round rule limits wealth effects, hedging across rounds, and excessive payment while preserving incentives in every decision.

3.3 Block 1: crisis and leader authorization

After consent, participants receive common instructions, complete baseline beliefs about public-service effectiveness and democratic risk, and answer a seven-item democratic-constraint battery. They then complete an unpaid guided practice covering citizen and leader decisions, the ballot, and random-round payment. Three comprehension questions must be answered correctly before production play begins; responses are optional in the development configurations. The Block-1 instructions describe a hypothetical society in a prolonged public-service crisis. When citizens decide individually, each point left in the fund produces 1.50 points for the group. When a leader decides, each point left in the fund produces 2.50 points, but the leader may move as many as 20 fund points to a personal account.

Participants then play ten paid rounds. Every round follows the same sequence: secret institutional ballot; majority rule; fund decisions under D or leader decisions under A; simultaneous feedback; and rematching. After Round 10, beliefs about public-service effectiveness and democratic risk and the constraint battery are measured again. The software records the final ballot, leader support in Rounds 8–10, all fund decisions, all personal transfers, and realized payoffs. One Block-1 round is randomly chosen for payment.

3.4 Block 2: structural change and reversal

Immediately after Block 1, each ten-person pool is assigned to one of two environments. Under *recovery*, the citizen-led multiplier rises from 1.50 to 2.50. Under *persistence*, it remains 1.50. The leader-led multiplier, authority, and maximum personal transfer remain fixed. A transition page explicitly states the rates that apply in Block 2. The treatment transparently equalizes productivity under recovery while preserving the leader-led method’s compulsory coordination.

Participants then play ten additional rounds. The same ballot, leader rule, group size, and payoff sequence make behavior directly comparable across blocks. After Round 10, the measures of public-service effectiveness and democratic risk and the constraint battery are repeated. The final Block-2 ballot is the reversal endpoint. One Block-2 round is randomly paid; the two selected-round payoffs are converted to euros and added to the participation payment.

4 Learning Across Repeated Play: Design Safeguards

Repeated play is an advantage only if procedural learning is not confused with the substantive learning of interest. The design therefore anticipates six implementation concerns.

Procedural learning. Participants may initially misunderstand payoff rules and improve merely through practice. A guided unpaid block and mandatory comprehension checks move rule learning before the paid data. Round 1 remains available for diagnostics, while the preregistered endpoint uses Round 10 and the late-round mean uses Rounds 8–10.

Reputation and reciprocity. Fixed five-person groups would allow personal histories to drive votes and contributions. Anonymous rematching within ten-person pools breaks stable bilateral relationships while keeping treatment exposure and synchronized interaction feasible. Participants never receive persistent identities.

Leader-role learning. Subjects could support delegation because they expect to become leader. Assignment is balanced among the least-served eligible members, preventing a few participants from monopolizing the role. Analyses will nevertheless describe prior leader service and realized personal transfers.

Endogenous institutional experience. Groups that never elect one procedure receive less experience with it. Instructions and practice expose everyone to both payoff rules, and realized experience is recorded. Treatment effects will be reported by prior exposure, but the primary intention-to-treat contrast will not condition on a post-treatment institutional path.

Wealth, fatigue, and end-game effects. Only one randomly selected round per block determines game earnings, so earnings from early rounds do not alter later endowments. Instructions explain that every round has an equal payment chance. The ten-round horizon is disclosed and identical across conditions. Page timers, decision times, and timeouts identify fatigue or end-game behavior. The persistence group experiences the same twenty-round sequence as the recovery group, separating the treatment effect from generic within-session fatigue.

Cross-block comparability. The method labels, leader rules, endowments, sequence, pool structure, and number of rounds are held constant. Recovery changes only the citizen-led fund multiplier, equalizing productivity while preserving the coordination asymmetry between methods. The persistence group identifies change caused by additional play, repeated measurement, or approaching the end of the session.

These safeguards do not attempt to eliminate all learning. The study requires participants to learn from consequences. They instead separate theoretically meaningful learning about institutional performance from avoidable learning about the interface or accidental partners.

5 Outcomes, Process Measures, and Estimands

Let $V_{ibr} \in \{D, A\}$ be participant i 's ballot in block b and round r . Define Block-1 leader endorsement as

$$A_{i1} = \mathbb{1}\{V_{i,1,10} = A\}. \quad (4)$$

For the theoretically relevant subgroup with $A_{i1} = 1$, democratic reversal is

$$R_i = \mathbb{1}\{V_{i,2,10} = D\}. \quad (5)$$

Let $T_g = 1$ indicate that the public-service crisis receded for matching pool g . The primary causal estimand is

$$\tau_R = \mathbb{E}[R_i(1) - R_i(0) \mid A_{i1} = 1], \quad (6)$$

estimated as an intention-to-treat contrast with standard errors clustered by the ten-person randomization pool. Because A_{i1} is measured in Block 1, before the Block-2 condition is disclosed, defining this subgroup does not condition on a post-treatment variable. The primary model will include prespecified baseline covariates only for precision; an unadjusted difference in means will also be reported.

In simple terms

The primary causal estimand (τ_R) asks: among participants who chose “A leader decides” in the final round of Block 1 ($V_{i,1,10} = \text{A}$, summarized by $A_{i1} = 1$), how much does assignment to recovery ($T_g = 1$), rather than persistence ($T_g = 0$), increase the probability that they return to “Citizens decide” in the final round of Block 2 ($V_{i,2,10} = \text{D}$), thereby constituting a democratic reversal ($R_i = 1$)? Under recovery, this return occurs after both methods reach the same productivity rate and while the leader-led method retains its ability to require equal allocations. The outcome therefore captures a consequential choice for collective control under continuing free-riding risk. Ideally, we would compare each participant’s decision under recovery ($R_i(1)$) with the decision that the same participant would have made under persistence ($R_i(0)$). Because both outcomes cannot be observed for the same person, random assignment allows us to estimate this effect by comparing reversal rates between conditions. Participants are analyzed according to their pool’s assigned condition (T_g), and statistical uncertainty is calculated at the pool level.

The ballot measures authorization of leader discretion; the transfer ρ_{igbr} measures its use. Define enacted appropriation among selected leaders as $M_{igbr} = \mathbb{1}\{\rho_{igbr} > 0\}$. Its intensity is recorded both in points and as a share of the available fund, $\rho_{igbr}/(5t_{igbr})$, when $t_{igbr} > 0$. These outcomes are individual choices, but leader selection and prior institutional outcomes are endogenous. They are therefore treated as theoretically important behavioral and process measures rather than as randomized effects of recovery.

Three secondary ballot outcomes describe robustness: (1) the change in leader-vote share from Block-1 Rounds 8–10 to Block-2 Rounds 8–10; (2) the full round-by-treatment trajectory; and (3) group-level selection of each method. Fund allocations, required allocations, personal transfers, fund size, leader service, exposure to appropriation, and realized payoffs describe the strategic environment in which ballots change.

Process outcomes occur at baseline, after Block 1, and after Block 2. They include perceived public-service effectiveness, perceived risk to constraints, the seven-item constraint index, round-level ballots, fund decisions, leader decisions, and realized payoffs. Battery items are randomized once per participant and then stored, preventing re-randomization on refresh. Any decomposition through perceived effectiveness or exposure to leader behavior will be labeled exploratory unless supported by a preregistered causal mediation strategy. The clean confirmatory claim remains the pool-randomized effect of recovery on reversal.

6 Randomization, Sample, and Power

The planned sample is 400 participants organized into 40 ten-person matching pools. For administrative reproducibility, pools are assigned in approximately equal numbers to recovery and persistence when the session is created. Assignment changes no Block-1 parameter and is disclosed only at the transition to Block 2. Eliminating the return visit removes cross-visit attrition and broken panel links. Recruitment should nevertheless overbook each synchronized cohort and maintain a standby list because interaction requires complete five-person groups at the start of the session.

Power depends on three quantities that cannot responsibly be inferred from the twenty repeated decisions: the proportion endorsing A at the Block-1 endpoint, the persistence-group reversal rate, and the intrapool correlation. A pilot will estimate these values, after which power will be evaluated by simulation at the pool level under the planned assignment. Repeated rounds increase

measurement reliability and identify trajectories; they do not turn 400 participants or 40 randomized pools into 8,000 independent observations.

Randomization code assigns whole ten-person pools and approximately balances the number of pools across conditions. Production sessions should contain a multiple of ten participants. If several laboratory rooms run concurrently, randomization and matching remain within the configured cohort so that participants never interact across inconsistent treatment states. **say it's better in-lab to avoid bot/AI bots.**

7 Payments and Laboratory Operations

At EUR 0.20 per point, the implemented parameters are designed to yield an expected total payment of roughly EUR 18–25, including the EUR 10 participation payment, depending on play. Final budgeting will use pilot distributions rather than theoretical maxima. At completion, the interface displays the selected Block-1 round, selected Block-2 round, point total, performance payment, participation payment, and total payment separately.

Both blocks are run during one synchronized session. The transition page is released only after every active participant has completed Block 1, so members of a matching pool encounter the new environment together. Participant codes, not names, link decisions across blocks, and all payments are made after Block 2.

The operational protocol requires complete groups. Standby participants may enter only before the first paid round. Once play begins, no replacement subject inherits another participant's history. Every strategic decision has a 90-second limit. If a participant times out, the software records the timeout and applies a neutral, prespecified default so that group members can continue. The session log must document late arrival, browser failure, early departure, manual intervention, and payment exceptions. Exclusion rules should be preregistered and should never depend on treatment-consistent behavior.

Before launch, two pilots are required. A technical pilot should use at least ten simultaneous clients and both treatment branches to stress wait pages, rematching, refresh persistence, timeouts, and payment calculations. A behavioral pilot should estimate duration, comprehension, Block-1 leader endorsement, leader appropriation, fatigue, and payoff distributions. Parameters should be frozen after that pilot and before confirmatory registration.

8 Ethics and Interpretation

The experiment exposes participants to a setting in which concentrated authority can be materially effective. That feature is necessary to study conditional tolerance, but it creates an obligation to avoid deception and moral labeling. Instructions therefore use the neutral labels “Citizens decide” and “A leader decides,” explain the public-services fund in plain language, and disclose exactly what the leader may do. The recovery condition is implemented through the programmed fund multiplier and explicitly announced before Block 2.

Consent explains the one-session structure, interaction, anonymity, payment rule, data linkage, and right to withdraw. The debrief explains the democratic interpretation only after all incentivized decisions and clarifies that no single vote or transfer establishes that a participant is corrupt, authoritarian, or opposed to democracy. Data exports should replace participant labels with study IDs and store the linkage key separately.

The Finnish setting strengthens the case for a hypothetical, incentive-based task. The experiment does not ask participants to judge whether Finnish officials appropriate public resources, nor does it

present the scenario as a description of Finland. Instead, all participants face the same fictional public-service crisis and make consequential choices under transparent rules. Variation in their behavior therefore speaks to individual willingness to preserve collective control when concentrated authority offers material benefits, even in a population for whom public appropriation may otherwise seem remote.

The interpretation centers on revealed institutional preference within a specified public-good environment. Recovery places both methods on the same productivity rate while preserving the leader’s power to coordinate contributions. An observed return to citizen control therefore shows how participants value democratic constraints when those constraints carry a continuing strategic cost. The randomized change in this trade-off and the within-person return to constraints form the study’s central contribution.

9 Planned Deliverables and Decision Rule

The confirmatory package will contain the theory, hypotheses, full instructions, oTree code, randomization procedure, exclusion rules, primary estimand, and pool-level power simulation. The principal conclusion will follow a simple rule: evidence for structural democratic reversal requires a positive, substantively meaningful, and statistically distinguishable recovery effect on R_i among Block-1 leader endorsers. Round-level trajectories, beliefs, fund decisions, leader choices, personal transfers, and allocations will be used to assess whether the observed endpoint follows the proposed learning mechanism and to identify rival explanations.

This design preserves the project’s core ambition—observing people learn and then potentially unlearn tolerance for weakly constrained authority—while giving that learning a strategic and distributive foundation. Subjects vote over rules that govern other subjects, may personally exercise the discretion those rules create, experience the consequences, and choose again after the crisis-created productivity premium for one-person rule expires. Democratic reversal is consequently revealed through a return to collective control while compulsory coordination remains economically attractive.

References

- Ananda, Aurelia and Damien Bol (2021). “Does Knowing Democracy Affect Answers to Democratic Support Questions? A Survey Experiment in Indonesia.” In: *International Journal of Public Opinion Research* 33.2, pp. 433–443.
- Capoccia, Giovanni (2007). *Defending Democracy*. Johns Hopkins University Press.
- Graham, Matthew and Milan Svolik (2020). “Democracy in America? Partisanship, Polarization, and the Robustness of Support for Democracy in the United States.” In: *American Political Science Review* 114.2, pp. 392–409.
- Grillo, Edoardo and Carlo Prato (2023). “Reference Points and Democratic Backsliding.” In: *American Journal of Political Science* 67.1, pp. 71–88.
- Hamman, John R., Roberto A. Weber, and Jonathan Woon (2011). “An Experimental Investigation of Electoral Delegation and the Provision of Public Goods.” In: *American Journal of Political Science* 55.4, pp. 738–752.
- Kingzette, Jonathan et al. (2021). “How Affective Polarization Undermines Support for Democratic Norms.” In: *Public Opinion Quarterly* 85.2, pp. 663–677.
- Levitsky, Steven and Daniel Ziblatt (2018). *How Democracies Die*. Crown.

- Saikkonen, Inga A.-L. and Henrik Serup Christensen (2023). “Guardians of Democracy or Passive Bystanders? A Conjoint Experiment on Elite Transgressions of Democratic Norms.” In: *Political Research Quarterly* 76.1, pp. 127–142.
- Wunsch, Natasha, Marc S. Jacob, and Laurenz Derksen (2025). “The Demand Side of Democratic Backsliding: How Divergent Understandings of Democracy Shape Political Choice.” In: *British Journal of Political Science*.